

Technical documentation

Contents

1 Lumen Geoguesser	1
1.1 Notices:	1
1.2 Directory Structure	1
1.3 Setup	1
1.3.1 Virtual Environment	1
1.3.2 Dataset Setup	2
1.4 Training	3
1.5 Logs - Tensorboard	4
1.6 Local Server	4

1 Lumen Geoguesser

1.1 Notices:

Although you might be reading this documentation in the form of a PDF file, **we strongly recommend that you open the `README.md` file in a markdown editor** (GitHub, VSCode, PyCharm, IDE...). As for the API documentation, after setting up the environment, we recommend you run the server with the `python3 src/app/main.py` command after which you can inspect API endpoints in a browser (and execute them too!). Essentially, the technical documentation PDF is rendered from the `README.md` markdown file and concatenated with the PDF API documentation.

A few more notes:

- the documentation assumes you are located at the `.lumen-geoguesser` directory when running Python scripts
- all global variables are defined in `src/config.py` and `src/paths.py`
- other directories have their own `README.md` files which hopefully will come in handy
- you can run most python files with the `python3 program.py -h` command to get a sense of which arguments you can/must send and what the script actually does

1.2 Directory Structure

Directory	Description
<code>data</code>	dataset, csvs, country shapefiles
<code>docs</code>	documentation
<code>figures</code>	figures
<code>models</code>	model checkpoints, model metadata
<code>references</code>	research papers and competition guidelines
<code>reports</code>	model checkpoints and model metadata
<code>src</code>	python source code

1.3 Setup

1.3.1 Virtual Environment

Create and populate the **virtual environment**. Simply put, the virtual environment allows you to install Python packages for this project only (which you can easily delete later). This way, we won't clutter your global Python packages.

Step 1: Execute the following command: - the command will initialize the `venv` if it doesn't exist yet

```
[ ! -d "venv" ] && (echo "Creating python3 virtual environment"; python3 -m venv venv)
. venv/bin/activate
pip install -r requirements.txt
```

1.3.2 Dataset Setup

This project allows for the usage multiple datasets, therefore, multiple dataset directories can usually be sent to `*.py` programs

Step 1: If needed, rename directory `data` (which contains uuid subdirectories) to `images` - The original dataset structure has a directory `data` (e.g `dataset_original_subset/data`) which contains subdirectories with uuids of locations (`dataset_original_subset/data/6bde8efe-a565-4f05-8c60-ae2ffb32ee9b`).

Dataset structure should look like this:

```
dataset_original_subset/
├── images
│   ├── 6bde8efe-a565-4f05-8c60-ae2ffb32ee9b
│   │   ├── 0.jpg
│   │   ├── 180.jpg
│   │   ├── 270.jpg
│   │   └── 90.jpg
│   ├── 6c0ed2ea-b31b-4cfd-9828-4aec22bc0b37
│   │   ├── 0.jpg
│   │   └── 180.jpg
│   └── ...
├── ...
└── data.csv
```

```
dataset_external_subset/
├── images
│   ├── e61b6e5f-db0d-4f57-bbe3-4d31f16c5bc3
│   │   ├── 0.jpg
│   │   └── 180.jpg
│   └── ...
├── ...
└── data.csv
```

Step 2: Setup datasets with `src/preprocess_setup_datasets.py` - Before running other scripts, you have to properly setup a new dataset structure using the `src/preprocess_setup_datasets.py` file. It's important to note that this file accepts multiple dataset directories as an argument and it will make sure to merge the datasets correctly. No changes will be done to your original directories.

```
python3 src/preprocess_setup_datasets.py -h
```

```
usage: preprocess_setup_datasets.py [-h] [--dataset-dirs dir [dir ...]] [--out-dir dir] [--copy-images] [--spacing SPACING]
```

optional arguments:

`-h, --help` show this help message and exit

`--dataset-dirs dir [dir ...]`

Dataset root directories that will be transformed into a single dataset

`--out-dir dir` Directory where the complete dataset will be placed

`--copy-images` Copy images from dataset directories to the new complete directory.

You don't need to do this as later on you will be able to pass multiple dataset directories to var

`--spacing SPACING`

Spacing that will be used to create a grid of polygons.

Different spacings generate a different number of classes

0.7 spacing => ~31 classes

0.5 spacing => ~55 classes

0.4 spacing => ~75 classes

0.3 spacing => ~115 classes

Example of running the initial setup script:

```
python3 src/preprocess_setup_datasets.py \  
--dataset-dirs data/dataset_original_subset data/dataset_external_subset \  
--out-dir data/dataset_complete_subset
```

`preprocess_setup_datasets.py` does all the necessary preprocessing. However, underneath the hood it calls other preprocessing scripts. What happens when you run this script?

1. a directory for the new (complete) dataset is created, images are copied if `--copy-images` flag was passed
2. `preprocess_csv_concat.main()` is called, which concatenates multiple `data.csv` s into a single `data.csv`
3. this new (complete) `data.csv` is enriched by `preprocess_csv_create_rich_static.main()`. Here, regions (future classes) and their information (centroids, crs centroids ...) are attached to each location. Enriched data is saved to a *Rich static CSV* file created called `data_rich_static_spacing_<float>_classes_<int>`.
4. Directory `images` in all directories (including the `complete` one) will be split into `train`, `val` and `test` directories.
Note: directory `images` won't be deleted.

New dataset structure:

```
dataset_complete_subset/  
├─ data.csv  
├─ images <= exists if the --copy-images flag was passed  
└─ data_rich_static_spacing_0.5_classes_55.csv  
  
dataset_original_subset/  
├─ data.csv  
├─ images  
├─ test  
│   └─ c4a74f0d-7f30-4966-9b92-f63279139d68  
│       ├── 0.jpg  
│       ├── 180.jpg  
│       └─ ...  
├─ ...  
├─ train  
└─ val  
  
dataset_external_subset/  
├─ data.csv  
├─ images  
├─ test  
├─ train  
└─ val
```

1.4 Training

After you prepared that new dataset structure, you can start the *quick version* of training:

```
python3 src/train.py --dataset-dirs data/dataset_external_subset/ data/dataset_original_subset/ \  
--csv-rich-static data/dataset_complete_subset/data_rich_static_spacing_0.7_classes_31.csv \  
--quick
```

You can stop the training anytime with `Ctrl + C`. Pressing it once will gracefully shutdown the training (and perform the testing phase). Pressing it twice shows more aggression, which will stop the training immediately.

`--csv-rich-static` can be left out which forces the *Rich static CSV* creation during runtime (this will somewhat slow down the initial setup). You can perform the full training by removing the `--quick` flag. Some additional interesting arguments are listed below. Run the `python src/train.py -h` command to see all supported arguments.

```
--image-size  
--num-workers  
--lr
```

```
--dataset-dirs [dir1, dir2, ...]
--csv-rich-static
--unfreeze-blocks
--pretrained
--quick
--batch-size
--optimizer
--regression
```

Example of production training (in our case):

```
python3 src/train.py \
--accelerator gpu --devices 1 --num-workers 32 \
--dataset-dir data/raw/ data/external/ \
--csv-rich-static data/complete/data_huge_spacing_0.21_num_class_211.csv \
--batch-size 8 --image-size 224 --lr 0.00002 \
--unfreeze-at-epoch 1 --scheduler plateau
```

During the training, a few things will occur in the `reports/` directory:

1. `reports/train_*.txt` files will be created which log everything that's outputted to the standard output
2. subdirectory `reports/<model_name>` will be created in which:
 1. `data_runtime.csv` will be created, serves as backup
 2. `version/0` directory which contains:
 1. `hparams.yaml` : details of hyperparameters
 2. `events.out.tfevents*` : log file which tensorboard consumes
 3. `checkpoints` : the most important subdirectory, contains model checkpoints (trained models)

```
reports/<model_name>/
├── data_runtime.csv
├── version_0
│   ├── checkpoints
│   │   ├── mymodel_checkpoint1.ckpt
│   │   └── mymodel_checkpoint2.ckpt
│   ├── events.out.tfevents.*
│   └── hparams.yaml
```

1.5 Logs - Tensorboard

Tensorboard logging is enabled by default. To see training and validation logs, run the command bellow. Logs should be available in a browser at `http://localhost:6006/` . For more options, check `tensorboard -h` .

```
tensorboard --port 6006 --logdir reports/
```



1.6 Local Server

A local server is useful when you are trying to do inference on a trained model. The sever code and config live in the `src/app` directory.

Before running the sever, set the variable `MODEL_DIRECTORY` in `src/app/.env` to a directory which contains (or will contain) model checkpoints (`.ckpt`). Models outside of this directory can't be used for inference via endpoints. We recommend creating a new directory called `models` and copying model checkpoint files (e.g. `reports/<model_name>/version_0/checkpoints/mymodel.ckpt`) to this directory.

Step 1. copy model checkpoints to /models/

```
mkdir models
cp -r reports/<model_name>/version_0/checkpoints/* models/
```

Step 1.1. ensure that the `MODEL_DIRECTORY` variable is set in `src/app/.env` file:

```
cat src/app/.env
```

Output:

```
MODEL_DIRECTORY = "models" # relative to the lumen-geoguesser directory
MODEL_EXTENSION = ".ckpt"
PREDICT_BATCH_SIZE = 16
HOST = "localhost"
PORT = 8090
HOT_RELOAD = 0 # please don't enable hot reloading as it's unstable
```

Step 2. run the server:

```
python3 src/app/main.py
```

FastAPI 0.1.0 OAS3
/openapi.json

Before consuming the endpoints set the variable `MODEL_DIRECTORY` in `.env` to a directory that contains model checkpoints (.ckpt). Models outside of this directory can't be used for inference via endpoints.

Command for running the server:

```
(venv) username@pc:~/lumen-geoguesser$ python3 src/app/main.py
```

available models All available models from the `MODEL_DIRECTORY` directory defined in the `.env` file ^

GET `/models` Get Models ✓

Endpoints used for predicting latitude and longitudes for given data. Curl example for multiple images:

predict ^

```
curl -i \
-F "images=@data/raw/images/train/e788b3d1-9d20-466c-9dee-97982f0f9a3b/0.jpg" \
-F "images=@data/raw/images/train/e788b3d1-9d20-466c-9dee-97982f0f9a3b/0.jpg" \
http://0.0.0.0:8090/model/Golf_76_haversine_0.0098_val_acc_0.47_val_loss_1.98_05-04-03-36-32/predict-images
```

POST `/model/{model_name}/predict-images` Predict Images ✓

POST `/model/{model_name}/predict-cardinal-images` Predict Cardinal Images ✓

POST `/model/{model_name}/predict-directory` Predict Dataset ✓

Schemas ^

Body_predict_cardinal_images_model__model_name__predict_cardinal_images_post >